

# Scalable Web Architectures(SWA)

## COMP 599 - Graduate Seminar, Fall 2018

Rihan Pereira, MSCS

Department of Computer Science  
California State University, Channel Islands

April 9, 2019



- These systems are designed to handle heavy web traffic, tolerate faults and yet continue to run

- These systems are designed to handle heavy web traffic, tolerate faults and yet continue to run
- Uses lot of research from Distributed Systems

- These systems are designed to handle heavy web traffic, tolerate faults and yet continue to run
- Uses lot of research from Distributed Systems
- SWA has a solid plan when failure happens - unreliable network, power outage, etc.

- These systems are designed to handle heavy web traffic, tolerate faults and yet continue to run
- Uses lot of research from Distributed Systems
- SWA has a solid plan when failure happens - unreliable network, power outage, etc.
- There is no right/wrong answer when building such systems; uses design principles to benchmark solutions to the system design problems it encounters.

- These systems are designed to handle heavy web traffic, tolerate faults and yet continue to run
  - Uses lot of research from Distributed Systems
  - SWA has a solid plan when failure happens - unreliable network, power outage, etc.
  - There is no right/wrong answer when building such systems; uses design principles to benchmark solutions to the system design problems it encounters.
- Availability
  - Partial Tolerance
  - Consistency
  - scalability
  - Maintenance

# Why I switched my topic

- At first, I had decided to present concepts in blockchain, I have barely scratched the surface of blockchain from technical standpoint.
- I am still learning, its huge, exciting.



# Why I switched my topic

- At first, I had decided to present concepts in blockchain, I have barely scratched the surface of blockchain from technical standpoint.
- I am still learning, its huge, exciting.
- Its still early, this space doesn't have a strong talent yet - smart people flock to study Machine Learning :)

# Why I switched my topic

- At first, I had decided to present concepts in blockchain, I have barely scratched the surface of blockchain from technical standpoint.
- I am still learning, its huge, exciting.
- Its still early, this space doesn't have a strong talent yet - smart people flock to study Machine Learning :)
- Much of the innovation is happening outside academia

# Why I switched my topic

- At first, I had decided to present concepts in blockchain, I have barely scratched the surface of blockchain from technical standpoint.
- I am still learning, its huge, exciting.
- Its still early, this space doesn't have a strong talent yet - smart people flock to study Machine Learning :)
- Much of the innovation is happening outside academia
- demand exceeds supply

- Thinking about scaling in advance before its a requirement makes systems unnecessarily complex without any benefit.

- Thinking about scaling in advance before its a requirement makes systems unnecessarily complex without any benefit.
- Although, some forethought into the design can save substantial time and resources in the future.


- Thinking about scaling in advance before its a requirement makes systems unnecessarily complex without any benefit.
- Although, some forethought into the design can save substantial time and resources in the future.
- As a software developer, especially backend engineer, will be dealing with Distributed Systems at some point in their career path

- Thinking about scaling in advance before its a requirement makes systems unnecessarily complex without any benefit.
- Although, some forethought into the design can save substantial time and resources in the future.
- As a software developer, especially backend engineer, will be dealing with Distributed Systems at some point in their career path
- A large population of backend engineer wont have the opportunity to build such systems from scratch, but you still need to tune knobs(configure) correctly, take advantage of features to achieve that sweet spot.

- Thinking about scaling in advance before its a requirement makes systems unnecessarily complex without any benefit.
- Although, some forethought into the design can save substantial time and resources in the future.
- As a software developer, especially backend engineer, will be dealing with Distributed Systems at some point in their career path
- A large population of backend engineer wont have the opportunity to build such systems from scratch, but you still need to tune knobs(configure) correctly, take advantage of features to achieve that sweet spot.




## Horizontal



- ability to redirect request to another nodes (in short, resiliency) \*
- load balancing
- network calls RPC/REST
- data consistency issues
- system scales proportional to variable data sets \*

## Vertical



- has a single point of failure
- N/A
- inter-process communication (IPC) \*
- consistent \*
- hardware limit


Imagine a simple image upload service using a central server



img/data\_access

Very soon, your appln is a hit!  
You start getting lot of traffic

Imagine a simple image upload service using a central server

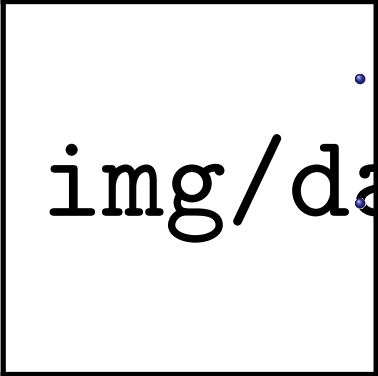


- Your service don't have high profit margins yet (bad reason, couldn't think of something else), so need to make sure the system design is cost-effective

img/data\_access

Very soon, your appln is a hit!  
You start getting lot of traffic

Imagine a simple image upload service using a central server

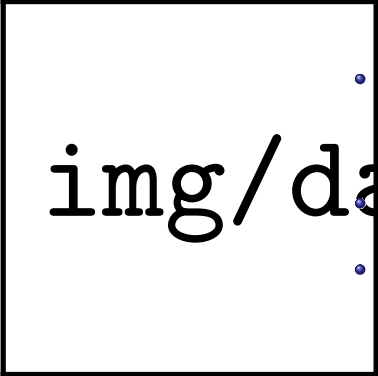


img/data\_access

- Your service don't have high profit margins yet(bad reason, couldn't think of something else), so need to make sure the system design is cost-effective
- Your system should be highly available, no zero tolerance

Very soon, your appln is a hit!  
You start getting lot of traffic

Imagine a simple image upload service using a central server

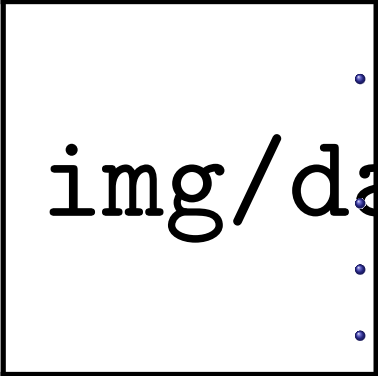


img/data\_access

- Your service don't have high profit margins yet(bad reason, couldn't think of something else), so need to make sure the system design is cost-effective
- Your system should be high available, no zero tolerance
- Upon image upload, image should be always there (you get what you put)

Very soon, your appln is a hit!  
You start getting lot of traffic

Imagine a simple image upload service using a central server



img/data\_access

- Your service don't have high profit margins yet(bad reason, couldn't think of something else), so need to make sure the system design is cost-effective
- Your system should be highly available, no zero tolerance
- Upon image upload, image should be always there (you get what you put)
- low-latency request-response

Very soon, your appln is a hit!  
You start getting lot of traffic

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)



## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)
- Availability

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)
- Availability
- Increased read throughput

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)
- Availability
- Increased read throughput

Challenge lies in handling changes to replicated data.

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)
- Availability
- Increased read throughput

Challenge lies in handling changes to replicated data.

- ① Single-leader based replication

## Replication

keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?

- keep data geographically close to your users (reduce latency)
- Availability
- Increased read throughput

Challenge lies in handling changes to replicated data.

- 1 Single-leader based replication
- 2 Multi-leader replication

## Replication


keeps a copy of the same data on multiple machines connected via a network.

Why you want to do it ?


- keep data geographically close to your users (reduce latency)
- Availability
- Increased read throughput

Challenge lies in handling changes to replicated data.

- 1 Single-leader based replication
- 2 Multi-leader replication
- 3 Leaderless replication



- Also called master-slave replication
- Considering synchronous replication, all nodes must be updated before the transaction is committed in the database.
- Using synchronous replication is not a good idea for high availability.
- Async configuration is widely used in production deployments.





# Follower setup without downtime

- Take a snapshot of the leader database at a specific time intervals

# Follower setup without downtime

- Take a snapshot of the leader database at a specific time intervals
- copy the snapshot to the new follower node.

# Follower setup without downtime

- Take a snapshot of the leader database at a specific time intervals
- copy the snapshot to the new follower node.
- followers connect to leaders & retrieves data changes happened since the snapshot was taken

- Take a snapshot of the leader database at a specific time intervals
- copy the snapshot to the new follower node.
- followers connect to leaders & retrieves data changes happened since the snapshot was taken
- when followers has processed the backlog of data changes since the last snapshot was taken, it is in sync.

- Take a snapshot of the leader database at a specific time intervals
- copy the snapshot to the new follower node.
- followers connect to leaders & retrieves data changes happened since the snapshot was taken
- when followers has processed the backlog of data changes since the last snapshot was taken, it is in sync.

How do you achieve high availability in this architecture ?

How do you achieve high availability in this architecture ?

- Follower failure: Catch-up recovery



How do you achieve high availability in this architecture ?


- Follower failure: Catch-up recovery
  - Each follower maintains a backlog


How do you achieve high availability in this architecture ?

- Follower failure: Catch-up recovery
  - Each follower maintains a backlog
- Leader failure

How do you achieve high availability in this architecture ?


- Follower failure: Catch-up recovery
  - Each follower maintains a backlog
- Leader failure
  - one of its followers takes the leader
  - clients need to be reconfigured to send writes to new leader
  - other followers need to start consuming data changes from new leader.
  - old leader joins back as normal follower






- natural extension of leader-based replica, allows more than 1 node to accept writes





- natural extension of leader-based replica, allows more than 1 node to accept writes
- each leader simultaneously accepts follow-up to each leader



- natural extension of leader-based replica, allows more than 1 node to accept writes
- each leader simultaneously accepts followers to each leader
- mostly used in multi-data center operations

# Conflict resolution problem




# Conflict resolution problem

- all replicas must arrive at the same final value.
- avoid conflict all-together


# Conflict resolution problem

- all replicas must arrive at the same final value.
- avoid conflict all-together
  - all writes for a particular record go through the same leader.
- In worst case, you have to deal with concurrent writes

- all replicas must arrive at the same final value.
- avoid conflict all-together
  - all writes for a particular record go through the same leader.
- In worst case, you have to deal with concurrent writes
- There is some research on conflict resolving
  - Conflict-free replicated data types
  - Mergeable persistent data structures
  - Operation transformation








- takes different approach instead of developing leader-follower concept






- takes different approach instead of developing leader-follower concept

- clients sends writes to






- takes different approach instead of developing leader-follower concept

- clients sends writes to

multiple nodes (peer-to-peer replication)

- No restriction on ordering






- takes different approach instead of developing leader-follower concept
- clients sends writes to
- No restriction on ordering
- Amazon's DynamoDB is built using this concept
- examples include Cassandra, Riak

- 3 replicas, 1 unavailable. how do you make up for unavailable replica ?

- 3 replicas, 1 unavailable. how do you make up for unavailable replica ?
- **Read Repair**
  - version 6 value from replica 3, version 7 value from replica 1 & 2 write version 7 to replica 3

- 3 replicas, 1 unavailable. how do you make up for unavailable replica ?
- **Read Repair**
  - version 6 value from replica 3, version 7 value from replica 1 & 2 write version 7 to replica 3
- **Anti-entropy process**
  - does reconciliation in background process
  - writes are copied in unordered way

Write processed on 2 out of 3 replicas. What if only 1 of 3 replicas accepted write ? How far can we push ?



Write processed on 2 out of 3 replicas. What if only 1 of 3 replicas accepted write ? How far can we push ?

Corollary


$$w + r > n$$



Write processed on 2 out of 3 replicas. What if only 1 of 3 replicas accepted write ? How far can we push ?

## Corollary

$$w + r > n$$



- If  $w < n$ , we can still process writes if a node is missing

- If  $r < n$ , we can still process reads if a node is missing

- $n = 3, w = 2, r = 2$ , we can tolerate 1 unavailable node
- $n = 5, w = 3, r = 3$ , we can tolerate 2 unavailable nodes

Very large datasets, having high query throughput are broken down into partitions or shards.





Very large datasets, having high query throughput are broken down into partitions or shards.

- think of each partition as a small database of its own
- approaches for partitioning large datasets
  - Partitioning by key range
  - Partitioning by hash of a key

Very large datasets, having high query throughput are broken down into partitions or shards.

- think of each partition as a small database of its own
- approaches for partitioning large datasets
  - Partitioning by key range
  - Partitioning by hash of a key





img

- a serious app will deploy a cache server
- usually placed in front of original data source
- algorithms like LRU (least recently used) are widely used
- problem is your system design uses a load balancer, you will need to overcome latency of the miles



figure 2




figure 3

1st figure handles eviction and data retrieval on its own

2nd figure application handles eviction

3rd figure each cache holds a portion of cache data

3rd figure uses consistent hashing to lookup data across nodes

basic role is to receive requests from client and relay them to next node in line.

basic role is to receive requests from client and relay them to next node in line.




figure 1

1st figure collapse similar requests into a single request




figure 2

2nd figure collapse requests for data that repeatedly lose together

A very common, popular and important technique used to speedup data access

A very common, popular and important technique used to speedup data access

img/indexing

- the trick is to index your database based on users data access patterns

img/indexlay

- In return for faster data access they do add write overhead and requiring updating indexes on each write



Their role is to distribute incoming requests evenly, fairly among available servers. They serve as a brokers between client and logical nodes, handling lot of simultaneous connections.

Their role is to distribute incoming requests evenly, fairly among available servers. They serve as a brokers between client and logical nodes, handling lot of simultaneous connections.



img/loadBalancer

Algorithms used -

- round robin
- just pick a random node
- selecting a node on a criteria - CPU, memory

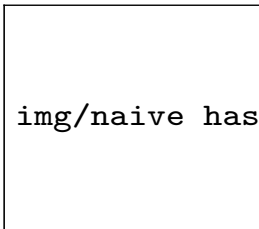
Queues solve a unique problem that load balancing, adding/removing servers cant solve.

Queues solve a unique problem that load balancing, adding/removing servers cant solve.

`img/queue_syn`

- work in async manner
- client is given acknowledgement that request is served
- tasks range from as simple as write to a data store to as complex as extracting pdf from text

`img/queue_async`




some assumptions -

- no. of nodes serving request never change
- repeated request can be served using cache
- requires rehashing every single key, caches get obsolete.

# Consistent Hashing

Incoming requests and serving nodes are placed onto a virtual ring structure called *hashring*



`img/consistent_has`

- placement of server nodes is not fixed on the ring, instead are placed at random locations
- each server owns a range of hashring
- No worries on adding new servers or server disruptions
- only rehashing of affected portion of requests is required



Martin Kleppmann,

*Designing Data Intensive applications, 2017*



James Hamilton,

*On Designing and Deploying Internet-Scale Services, 2007.*



Twitter,

*Automatic Management of Partitioned, Replicated Search Services, 2011.*



<http://blog.acolyer.org/>,

*The morning Paper.*



<http://aosabook.org>,

*The architecture of open source applications.*



<http://book.mixu.net/distsys>,

*Distributed Systems: fun and profit.*

Thank you! Questions ?